

Import Library

In [1]:

```
import pandas as pd
```

In [2]:

```
import numpy as np
```

In [3]:

```
import matplotlib.pyplot as plt
```

In [4]:

```
import seaborn as sns
```

Import Data

In [7]:

```
df = pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/MPG.csv')
```

In [8]:

```
df.head()
```

Out[8]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin	name
0	18.0	8	307.0	130.0	3504	12.0	70	usa	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693	11.5	70	usa	buick skylark 32
2	18.0	8	318.0	150.0	3436	11.0	70	usa	plymouth satellite
3	16.0	8	304.0	150.0	3433	12.0	70	usa	american rebel sports
4	17.0	8	302.0	140.0	3449	10.5	70	usa	ford torino



In [9]:

```
df.nunique()
```

Out[9]:

```
mpg          129
cylinders      5
displacement  82
horsepower    93
weight       351
acceleration  95
model_year    13
origin         3
name         305
dtype: int64
```

Data Preprocessing

In [11]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   mpg             398 non-null   float64
 1   cylinders       398 non-null   int64   
 2   displacement    398 non-null   float64
 3   horsepower      392 non-null   float64
 4   weight         398 non-null   int64   
 5   acceleration    398 non-null   float64
 6   model_year      398 non-null   int64   
 7   origin          398 non-null   object  
 8   name            398 non-null   object  
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

In [13]:

df.describe

Out[13]:

```
<bound method NDFrame.describe of
owner  weight  acceleration  \
0      18.0      8          307.0      130.0      3504      12.0
1      15.0      8          350.0      165.0      3693      11.5
2      18.0      8          318.0      150.0      3436      11.0
3      16.0      8          304.0      150.0      3433      12.0
4      17.0      8          302.0      140.0      3449      10.5
..      ...      ...          ...      ...      ...      ...
393    27.0      4          140.0      86.0      2790      15.6
394    44.0      4           97.0      52.0      2130      24.6
395    32.0      4          135.0      84.0      2295      11.6
396    28.0      4          120.0      79.0      2625      18.6
397    31.0      4          119.0      82.0      2720      19.4

      model_year  origin      name
0              70     usa  chevrolet chevelle malibu
1              70     usa      buick skylark 320
2              70     usa  plymouth satellite
3              70     usa      amc rebel sst
4              70     usa      ford torino
..            ...     ...      ...
393             82     usa  ford mustang gl
394             82  europe      vw pickup
395             82     usa  dodge rampage
396             82     usa  ford ranger
397             82     usa  chevy s-10
```

[398 rows x 9 columns]>

In [15]:

df.corr()

Out[15]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_y
mpg	1.000000	-0.775396	-0.804203	-0.778427	-0.831741	0.420289	0.579267
cylinders	-0.775396	1.000000	0.950721	0.842983	0.896017	-0.505419	-0.348746
displacement	-0.804203	0.950721	1.000000	0.897257	0.932824	-0.543684	-0.370164
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	-0.689196	-0.416361
weight	-0.831741	0.896017	0.932824	0.864538	1.000000	-0.417457	-0.306564
acceleration	0.420289	-0.505419	-0.543684	-0.689196	-0.417457	1.000000	0.288137
model_year	0.579267	-0.348746	-0.370164	-0.416361	-0.306564	0.288137	1.000000

=

Remove missing values

In [17]:

```
df=df.dropna()
```

In [19]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 392 entries, 0 to 397
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   mpg             392 non-null   float64
 1   cylinders       392 non-null   int64   
 2   displacement    392 non-null   float64
 3   horsepower      392 non-null   float64
 4   weight          392 non-null   int64   
 5   acceleration    392 non-null   float64
 6   model_year      392 non-null   int64   
 7   origin          392 non-null   object  
 8   name            392 non-null   object  
dtypes: float64(4), int64(3), object(2)
memory usage: 30.6+ KB
```

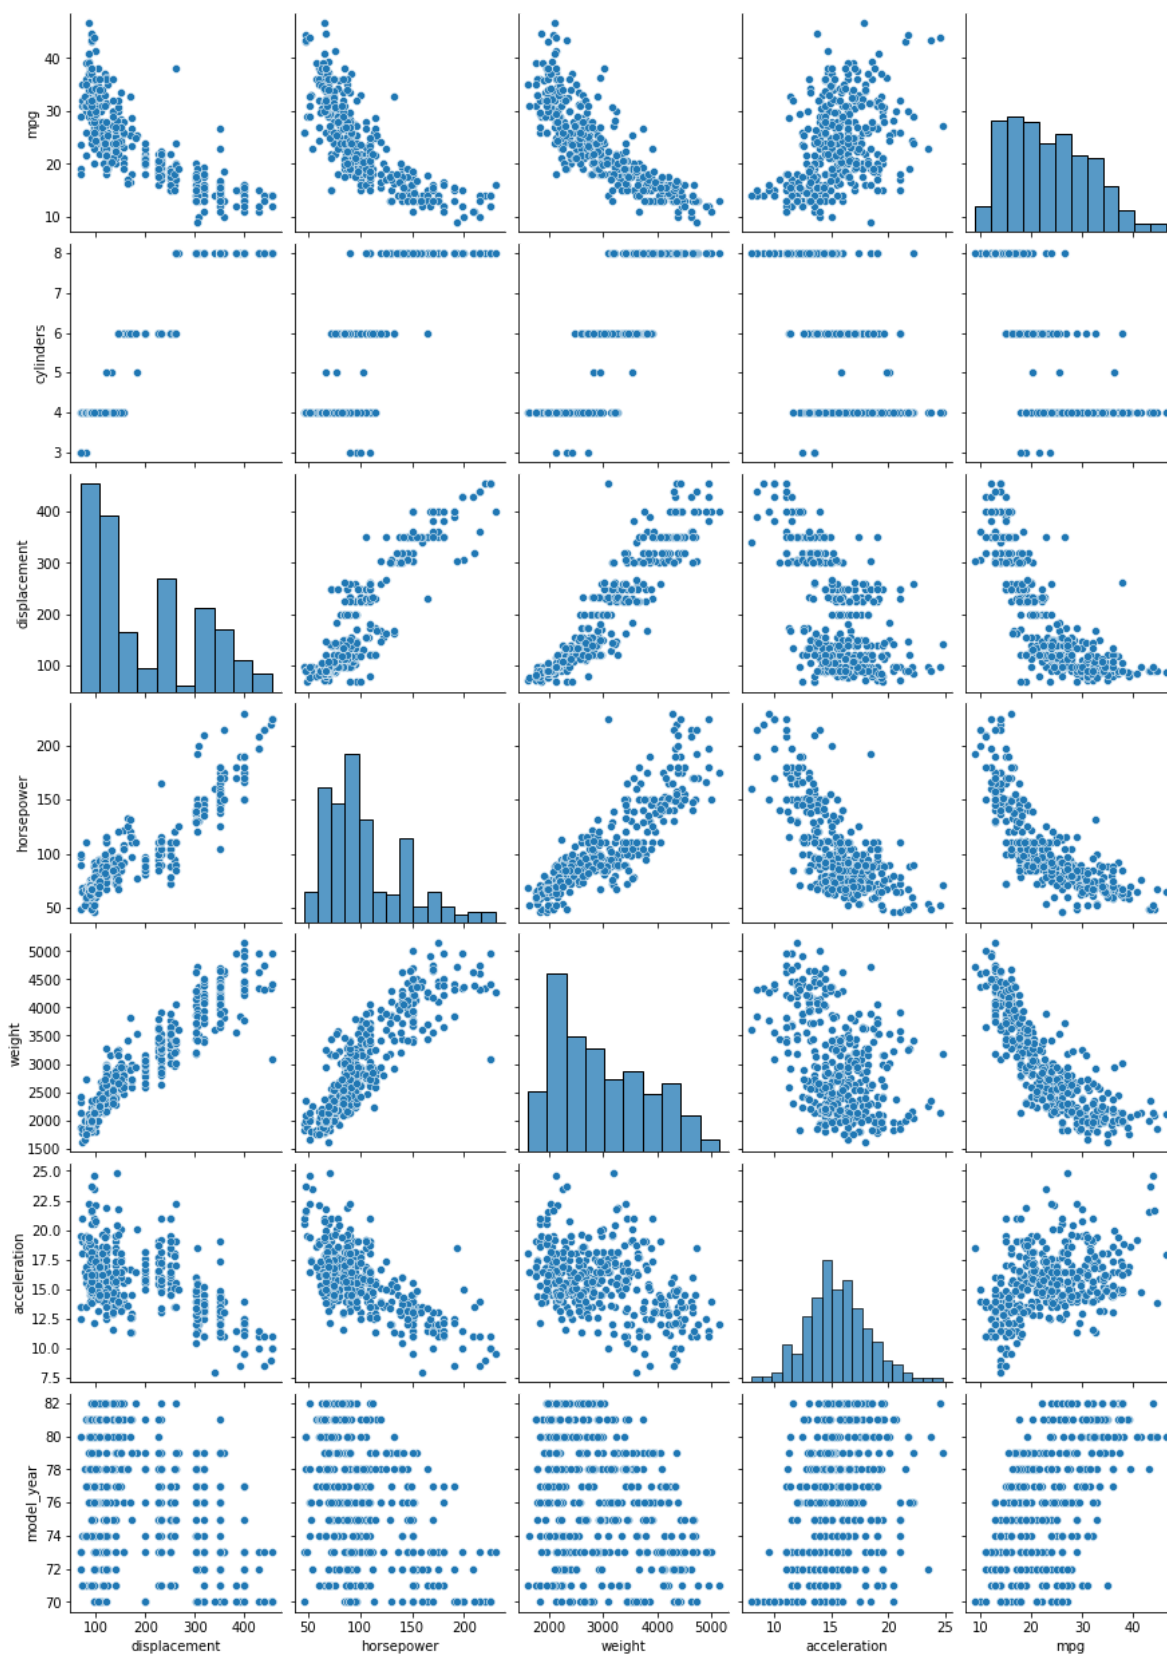
Data Visualization

In [23]:

```
sns.pairplot(df, x_vars=['displacement', 'horsepower', 'weight', 'acceleration', 'mpg'])
```


Out[23]:

<seaborn.axisgrid.PairGrid at 0x7f1d93e41450>



In [25]:

```
sns.regplot(x='displacement', y='mpg', data=df);
```



Define Target variable y and feature x

In [26]:

```
df.columns
```

Out[26]:

```
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',  
      'acceleration', 'model_year', 'origin', 'name'],  
      dtype='object')
```

In [28]:

```
y=df['mpg']
```

In [29]:

```
y.shape
```

Out[29]:

```
(392,)
```

In [67]:

```
X=df[['displacement', 'horsepower', 'weight', 'acceleration']]
```

In [68]:

```
X.shape
```

Out[68]:

```
(392, 4)
```

In [69]:

X

Out[69]:

	displacement	horsepower	weight	acceleration
0	307.0	130.0	3504	12.0
1	350.0	165.0	3693	11.5
2	318.0	150.0	3436	11.0
3	304.0	150.0	3433	12.0
4	302.0	140.0	3449	10.5
...
393	140.0	86.0	2790	15.6
394	97.0	52.0	2130	24.6
395	135.0	84.0	2295	11.6
396	120.0	79.0	2625	18.6
397	119.0	82.0	2720	19.4

392 rows × 4 columns



In [70]:

X.describe()

Out[70]:

	displacement	horsepower	weight	acceleration
count	392.000000	392.000000	392.000000	392.000000
mean	194.411990	104.469388	2977.584184	15.541327
std	104.644004	38.491160	849.402560	2.758864
min	68.000000	46.000000	1613.000000	8.000000
25%	105.000000	75.000000	2225.250000	13.775000
50%	151.000000	93.500000	2803.500000	15.500000
75%	275.750000	126.000000	3614.750000	17.025000
max	455.000000	230.000000	5140.000000	24.800000



Scaling Data

In [35]:

```
from sklearn.preprocessing import StandardScaler
```

In [36]:

```
ss=StandardScaler()
```

In [38]:

```
X=ss.fit_transform(X)
```

In [40]:

```
X
```

Out[40]:

```
array([[ 1.07728956,  0.66413273,  0.62054034, -1.285258  ],
       [ 1.48873169,  1.57459447,  0.84333403, -1.46672362],
       [ 1.1825422 ,  1.18439658,  0.54038176, -1.64818924],
       ...,
       [-0.56847897, -0.53247413, -0.80463202, -1.4304305 ],
       [-0.7120053 , -0.66254009, -0.41562716,  1.11008813],
       [-0.72157372, -0.58450051, -0.30364091,  1.40043312]])
```

In [41]:

```
pd.DataFrame(X).describe()
```

Out[41]:

	0	1	2	3
count	3.920000e+02	3.920000e+02	3.920000e+02	3.920000e+02
mean	-1.393443e-16	-3.293850e-16	5.607759e-17	1.608691e-16
std	1.001278e+00	1.001278e+00	1.001278e+00	1.001278e+00
min	-1.209563e+00	-1.520975e+00	-1.608575e+00	-2.736983e+00
25%	-8.555316e-01	-7.665929e-01	-8.868535e-01	-6.410551e-01
50%	-4.153842e-01	-2.853488e-01	-2.052109e-01	-1.499869e-02
75%	7.782764e-01	5.600800e-01	7.510927e-01	5.384714e-01
max	2.493416e+00	3.265452e+00	2.549061e+00	3.360262e+00

```
==
```

Train test split data

In [43]:

```
from sklearn.model_selection import train_test_split
```

In [45]:

```
X_train, X_test, y_train, y_test=train_test_split(X,y, train_size=0.7, random_state=2525)
```

In [49]:

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[49]:

```
((274, 4), (118, 4), (274,), (118,))
```


Linear Regression Model

In [50]:

```
from sklearn.linear_model import LinearRegression
```

In [52]:

```
lr=LinearRegression()
```

In [53]:

```
lr.fit(X_train, y_train)
```

Out[53]:

```
LinearRegression()
```

In [54]:

```
lr.intercept_
```

Out[54]:

```
23.688921610685803
```

In [56]:

```
lr.coef_
```

Out[56]:

```
array([-0.13510042, -1.4297211 , -5.23891463,  0.22436094])
```

Predict test data

In [57]:

```
y_pred=lr.predict(X_test)
```

In [59]:

```
y_pred
```

Out[59]:

```
array([25.24954801, 26.85525431, 26.58882904, 29.48052754, 23.91216916,
       14.9529791 , 30.0607685 , 34.07634195, 30.550342 , 11.31024173,
       18.14067535, 18.75305197, 29.80678264, 33.19954312, 17.23635872,
       16.06983768, 25.94812038, 21.15777548, 29.92508087, 25.05587641,
       22.85575427, 30.96630956, 22.82202336, 24.04513247, 25.95102384,
       26.21136844, 14.91805111, 31.85928917, 21.95227216, 26.85446824,
        8.94214825, 26.21244694, 30.20552304,  7.15733458, 26.31771126,
       30.54356872, 14.13603243, 31.02810818, 33.19140036, 31.74995879,
       11.07428823, 30.50398808, 29.36195486, 31.022648 , 23.53384962,
       22.87821543, 11.03531446, 14.3757476 , 31.44484893, 26.64255441,
       27.96470623, 21.80486111, 20.32272978, 31.27632871, 24.83127389,
       19.13391479, 28.2786737 , 25.21468804, 26.89045676, 28.76603057,
       19.03600671, 29.49310219, 28.42147856, 26.6112997 ,  7.384747 ,
       20.13152225, 22.77931428, 20.50765035, 32.81875326, 27.92430623,
       13.34341223,  8.03767139, 25.34229398, 17.23635872, 33.03710336,
       31.07878627, 21.58700058, 24.53266643, 30.38829664, 17.84737111,
       31.30622407, 30.1021144 , 22.81248978, 20.01904445,  9.12644754,
       24.50457451, 29.57695629, 29.45235437, 31.59169567, 26.49442535,
       30.32795983, 12.36145993, 16.48933189, 15.27329229, 32.77989962,
       27.25863029, 11.07878871, 25.72147567, 12.57968624, 30.4363069 ,
       27.56306784, 24.92600083, 16.21791725, 23.89776551, 18.63499966,
       10.21748386, 21.60970196, 23.01257072, 27.30850629, 30.45961552,
       29.43254102, 27.21176721, 24.2365775 , 28.87030773, 21.16703179,
       27.97152628, 24.54560958, 32.23487944])
```

Model Accuracy

In [71]:

```
from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error, r2_score
```

In [72]:

```
mean_absolute_error(y_test, y_pred)
```

Out[72]:

```
3.417654680078564
```

In [73]:

```
mean_absolute_percentage_error(y_test, y_pred)
```

Out[73]:

```
0.16282215595698368
```

In [74]:

```
r2_score(y_test, y_pred)
```

Out[74]:

0.6767436309121445

Polynomial Regression

In [75]:

```
from sklearn.preprocessing import PolynomialFeatures
```

In [78]:

```
poly=PolynomialFeatures(degree=2, interaction_only=True, include_bias=False)
```

In [79]:

```
X_train2=poly.fit_transform(X_train)
```

In [80]:

```
X_test2=poly.fit_transform(X_test)
```

In [82]:

```
lr.fit(X_train2, y_train)
```

Out[82]:

LinearRegression()

In [84]:

```
lr.intercept_
```

Out[84]:

21.457120355191684

In [85]:

```
lr.coef_
```

Out[85]:

```
array([-1.97594907e+00, -5.50639326e+00, -1.82341405e+00, -8.04049934e-01,  
       1.55534517e+00, -4.40583099e-01, -5.33735335e-01,  1.29466895e+00,  
       2.61553723e-03,  5.86761939e-01])
```

In [87]:

```
y_pred_poly=lr.predict(X_test2)
```

Model accuracy

In [89]:

```
from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error, r2_score
```

In [90]:

```
mean_absolute_error(y_test, y_pred_poly)
```

Out[90]:

2.924007242447459

In [92]:

```
mean_absolute_percentage_error(y_test, y_pred_poly)
```

Out[92]:

0.12874881331071994

In [93]:

```
r2_score(y_test, y_pred_poly)
```

Out[93]:

0.7198303534964864

Hand Written Digit Prediction - Classification Analysis

Import Library

In [94]:

```
import pandas as pd
```

In [95]:

```
import numpy as np
```

In [98]:

```
import matplotlib.pyplot as plt
```

Import Data

In [99]:


```
from sklearn.datasets import load_digits
```

In [100]:

```
df=load_digits()
```

In [102]:

```
_, axes=plt.subplots(nrows=1, ncols=4, figsize=(10,3))
for ax, image, label in zip(axes,df.images, df.target):
    ax.set_axis_off()
    ax.imshow(image, cmap=plt.cm.gray_r, interpolation="nearest")
    ax.set_title("Training: %i" % label)
```



Data Preprocessing

In [103]:

```
df.images.shape
```

Out[103]:

```
(1797, 8, 8)
```

In [104]:

```
df.images[0]
```

Out[104]:

```
array([[ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.],
       [ 0.,  0., 13., 15., 10., 15.,  5.,  0.],
       [ 0.,  3., 15.,  2.,  0., 11.,  8.,  0.],
       [ 0.,  4., 12.,  0.,  0.,  8.,  8.,  0.],
       [ 0.,  5.,  8.,  0.,  0.,  9.,  8.,  0.],
       [ 0.,  4., 11.,  0.,  1., 12.,  7.,  0.],
       [ 0.,  2., 14.,  5., 10., 12.,  0.,  0.],
       [ 0.,  0.,  6., 13., 10.,  0.,  0.,  0.]])
```

In [105]:

```
df.images[0].shape
```

Out[105]:

```
(8, 8)
```

In [107]:

```
n_samples=len(df.images)
data=df.images.reshape((n_samples, -1))
```

In [109]:

```
data[0]
```

Out[109]:

```
array([ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
        15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
        12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
         0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
        10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.]
```

In [110]:

```
data[0].shape
```

Out[110]:

```
(64,)
```

In [112]:

```
data.shape
```

Out[112]:

```
(1797, 64)
```

Scaling Data

In [113]:

```
data.min()
```

Out[113]:

```
0.0
```

In [114]:

```
data.max()
```

Out[114]:

```
16.0
```

In [116]:

```
data=data/16
```

In [117]:

```
data.min()
```

Out[117]:

```
0.0
```

In [119]:

```
data.max()
```

Out[119]:

0.0625

In [121]:

```
data[0]
```

Out[121]:

```
array([0.          , 0.          , 0.01953125, 0.05078125, 0.03515625,
        0.00390625, 0.          , 0.          , 0.          , 0.          ,
        0.05078125, 0.05859375, 0.0390625 , 0.05859375, 0.01953125,
        0.          , 0.          , 0.01171875, 0.05859375, 0.0078125 ,
        0.          , 0.04296875, 0.03125   , 0.          , 0.          ,
        0.015625 , 0.046875  , 0.          , 0.          , 0.03125   ,
        0.03125   , 0.          , 0.          , 0.01953125, 0.03125   ,
        0.          , 0.          , 0.03515625, 0.03125   , 0.          ,
        0.          , 0.015625 , 0.04296875, 0.          , 0.00390625,
        0.046875 , 0.02734375, 0.          , 0.          , 0.0078125 ,
        0.0546875 , 0.01953125, 0.0390625 , 0.046875  , 0.          ,
        0.          , 0.          , 0.          , 0.0234375 , 0.05078125,
        0.0390625 , 0.          , 0.          , 0.          ])
```

Train test split data

In [122]:

```
from sklearn.model_selection import train_test_split
```

In [127]:

```
X_train, X_test, y_train, y_test=train_test_split(data, df.target, test_size=0.3)
```

In [128]:

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[128]:

```
((1257, 64), (540, 64), (1257,), (540,))
```

Random Forest Model

In [130]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [132]:

```
rf=RandomForestClassifier()
```

In [133]:

```
rf.fit(X_train, y_train)
```

Out[133]:

```
RandomForestClassifier()
```

Predict Test Data

In [134]:

```
y_pred=rf.predict(X_test)
```

In [135]:

```
y_pred
```

Out[135]:

```
array([7, 4, 1, 3, 1, 5, 9, 3, 4, 7, 6, 4, 4, 8, 8, 2, 6, 0, 3, 7, 2, 8,  
       0, 5, 9, 6, 5, 1, 2, 0, 0, 5, 3, 7, 0, 2, 4, 7, 4, 3, 0, 1, 8, 2,  
       2, 7, 9, 9, 0, 6, 1, 4, 6, 3, 8, 8, 3, 3, 8, 9, 1, 3, 6, 1, 0, 1,  
       0, 2, 6, 6, 5, 2, 5, 7, 9, 9, 9, 4, 2, 3, 3, 5, 6, 5, 7, 9, 3, 4,  
       1, 7, 4, 3, 5, 8, 1, 0, 8, 2, 6, 1, 2, 6, 6, 1, 3, 6, 7, 4, 9, 1,  
       9, 9, 0, 1, 5, 8, 9, 2, 8, 8, 3, 8, 2, 8, 9, 1, 6, 0, 2, 3, 1, 9,  
       6, 3, 2, 8, 4, 2, 4, 7, 4, 0, 3, 2, 5, 5, 6, 1, 8, 3, 9, 9, 7, 2,  
       5, 7, 7, 1, 8, 8, 5, 5, 9, 5, 8, 9, 2, 1, 0, 6, 0, 8, 3, 4, 3, 5,  
       9, 5, 4, 3, 0, 4, 1, 0, 2, 6, 7, 1, 1, 6, 0, 5, 5, 9, 4, 0, 2, 5,  
       7, 7, 4, 0, 7, 4, 3, 0, 8, 5, 2, 3, 4, 2, 2, 3, 9, 0, 6, 7, 1, 0,  
       2, 7, 1, 1, 4, 1, 9, 8, 1, 9, 0, 6, 7, 6, 0, 1, 6, 9, 8, 7, 7, 2,  
       5, 2, 9, 2, 9, 2, 6, 9, 0, 9, 0, 1, 8, 4, 8, 1, 4, 2, 5, 1, 5, 0,  
       6, 1, 5, 5, 1, 6, 6, 3, 0, 8, 1, 1, 1, 1, 1, 1, 7, 1, 0, 3, 2, 0,  
       6, 4, 6, 9, 7, 2, 9, 6, 4, 9, 7, 3, 0, 3, 3, 2, 3, 3, 6, 6, 2, 4,  
       9, 1, 2, 1, 1, 3, 4, 3, 7, 5, 6, 0, 1, 5, 1, 1, 6, 5, 0, 3, 4, 8,  
       1, 3, 2, 2, 1, 0, 3, 8, 0, 6, 4, 2, 4, 0, 9, 5, 8, 9, 5, 9, 6, 6,  
       0, 6, 2, 2, 9, 7, 5, 5, 7, 8, 7, 5, 5, 4, 1, 7, 2, 5, 0, 7, 9, 1,  
       6, 4, 7, 1, 1, 2, 6, 9, 0, 8, 8, 7, 9, 6, 4, 3, 1, 9, 3, 4, 2, 6,  
       7, 9, 1, 6, 8, 8, 2, 5, 9, 7, 0, 1, 2, 5, 3, 8, 0, 8, 6, 6, 5, 9,  
       5, 7, 0, 5, 5, 9, 8, 8, 1, 2, 5, 1, 8, 4, 6, 7, 0, 2, 8, 8, 9, 4,  
       2, 4, 3, 4, 2, 3, 6, 9, 9, 8, 4, 4, 9, 6, 6, 4, 1, 6, 7, 9, 5, 5,  
       5, 9, 8, 0, 4, 5, 0, 5, 4, 1, 1, 3, 2, 2, 7, 8, 2, 8, 2, 2, 9, 1,  
       7, 4, 8, 7, 2, 1, 4, 9, 7, 5, 2, 8, 4, 1, 8, 3, 6, 3, 3, 6, 7, 0,  
       7, 1, 7, 3, 2, 3, 0, 2, 5, 7, 7, 6, 5, 9, 4, 1, 8, 2, 3, 1, 8, 4,  
       0, 0, 1, 7, 6, 7, 5, 9, 1, 3, 4, 0])
```

Model Accuracy

In [136]:

```
from sklearn.metrics import confusion_matrix, classification_report
```


In [137]:

```
confusion_matrix(y_test, y_pred)
```

Out[137]:

```
array([[51,  0,  0,  0,  1,  0,  0,  0,  0,  0],
       [ 0, 66,  0,  0,  0,  0,  0,  0,  0,  0],
       [ 0,  0, 59,  0,  0,  0,  0,  0,  0,  0],
       [ 0,  0,  0, 49,  0,  0,  0,  2,  0,  0],
       [ 0,  0,  0,  0, 48,  0,  0,  1,  0,  0],
       [ 0,  0,  0,  0,  0, 53,  0,  0,  0,  1],
       [ 0,  0,  0,  0,  0,  0, 54,  0,  0,  0],
       [ 0,  0,  0,  0,  0,  0,  0, 47,  0,  0],
       [ 0,  2,  0,  0,  1,  0,  0,  0, 49,  0],
       [ 0,  0,  0,  1,  0,  0,  0,  0,  1, 54]])
```

In [138]:

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	52
1	0.97	1.00	0.99	66
2	1.00	1.00	1.00	59
3	0.98	0.96	0.97	51
4	0.96	0.98	0.97	49
5	1.00	0.98	0.99	54
6	1.00	1.00	1.00	54
7	0.94	1.00	0.97	47
8	0.98	0.94	0.96	52
9	0.98	0.96	0.97	56
accuracy			0.98	540
macro avg	0.98	0.98	0.98	540
weighted avg	0.98	0.98	0.98	540